

Probability Estimation Using Monte Carlo Simulation of Boolean Logic on Hardware-Accelerated Platforms

Arjun Earthperson
PhD Candidate, PRA Group



NC STATE
UNIVERSITY

EMBEDDED TOPICAL CONFERENCE AT THE
 **ANS**® 2025 ANNUAL
CONFERENCE



└ Motivation

└ Evolving Hardware Landscape

Industry Response to Emergent AI/ML Workloads

Heavy Investment in Data-Parallel Hardware

- GPUs, tensor cores provide high throughput for integer operations.
 - Apple M4 Neural Engine: ≈ 38 TOPS (int8).
 - Nvidia RTX 4090: ≈ 1000 TOPS.
- Current-gen consumer hardware already supports specialized ops (Intel AMX).

Designed for Inference on Massive Models

- $\approx 10^9$ parameters on mobile devices (e.g., Gemma 2B).
- $\approx 10^{12}$ parameters on HPC/cloud (LLaMa 4 at 2 trillion).

Comparatively,

- Largest (public) PRA models: $\approx 10^6$ parameters (Generic PWR).
- However, PRA model quantification is deeply recursive.



Overview and Highlights I

High-Throughput Boolean Logic Evaluation (aka Eval Query)

- Simultaneous evaluation of *all* intermediate gates, success, and failure paths.
- Relax coherence constraints: arbitrary graph shapes, with NOT, or any other combination of gates permitted.
- Boolean logic manipulation not strictly needed.
- Vectorized bitwise hardware ops for logical primitives (AND, OR, XOR, etc.)
- Specialized treatment of k/n logic, without expansion.
- Concurrent use of all available compute - GPUs, multicore CPUs.



Overview and Highlights II

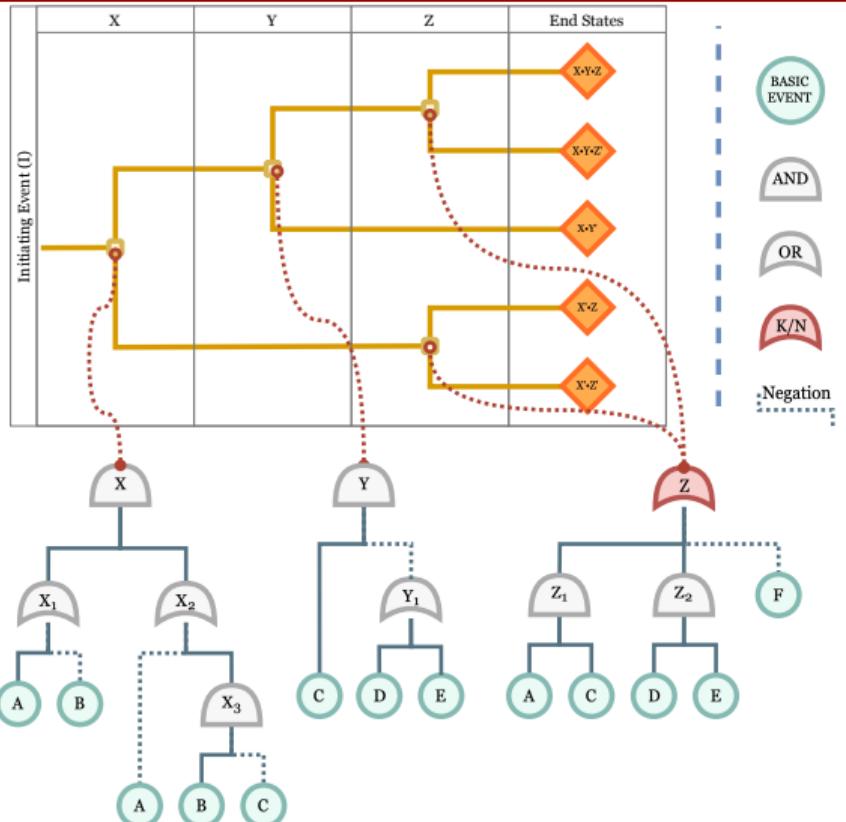
Probability Estimation using Eval Query (aka Expectation Query)

- Estimate probabilities for *all* events in the PRA model simultaneously.
- No need to compute the minimal cut sets or prime implicants.
- Streamable: Solve the entire PRA model iteratively, approx 0.3 seconds per iteration, regardless of model size. More iterations needed for complex models.



└ PRA Models as Probabilistic Circuits

└ A Working Example: One Initiating Event, Three Fault Trees, Six Basic Events, Five End States



Variable	Expression
X	$(A B') \bullet (A' (B \bullet C'))$
Y	$C \bullet (D E)'$
Z	$(A \bullet C) (D \bullet E) F'$

Table: Unimplified Boolean expression for each Top Event

Small, but non-trivial structure:

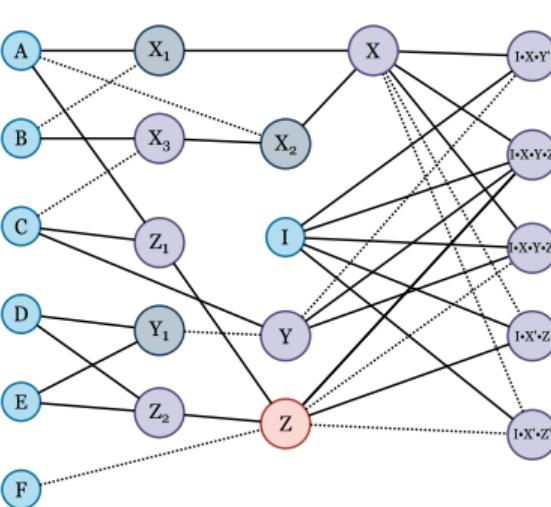
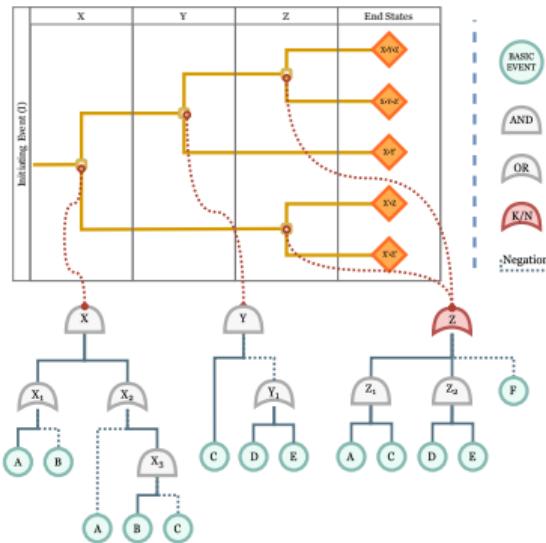
- Basic events are shared.
- Some gate outputs are negated.
- Event Z is a (k=2) of n=3 gate.



└ PRA Models as Probabilistic Circuits

└ A Working Example: One Initiating Event, Three Fault Trees, Six Basic Events, Five End States

Compile a Directed Acyclic Graph (DAG) from Logic Model



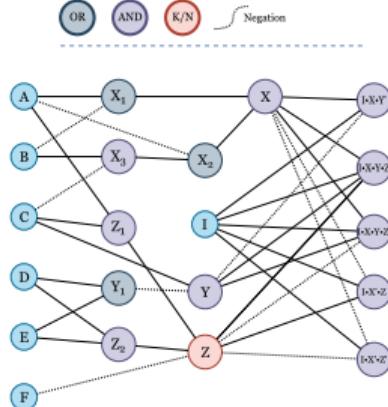
Start with an Arbitrary Topological Ordering:

- Aim for succinctness.
- prefer HW-native ops.
- e.g. don't expand k/n.



└ PRA Models as Probabilistic Circuits

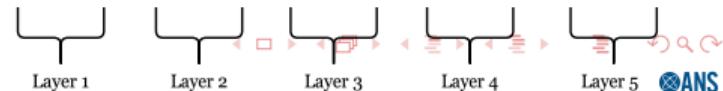
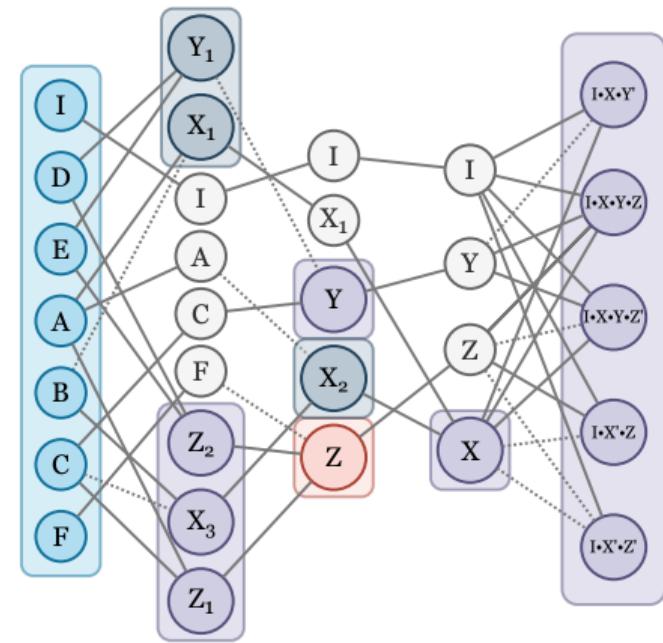
└ A Working Example: One Initiating Event, Three Fault Trees, Six Basic Events, Five End States



Refinements:

- Partition into layers.
- Vectorize for SIMD by fusing similar ops.
- Feed-forward only: improves caching.

Still not probabilistic: inputs, outputs are bitpacked INT64 tensors.





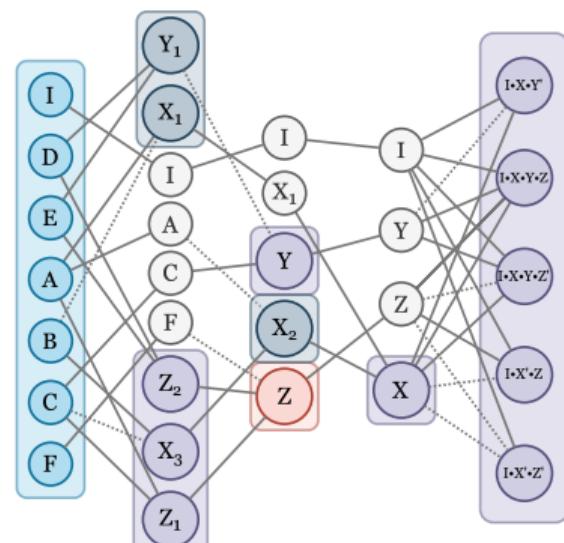
- └ PRA Models as Probabilistic Circuits
 - └ Knowledge Compilation and Queries

Querying the Compiled Knowledge Graph

The Simplest Type of Query: $\text{Eval}(G)$

- Set the inputs [on/off].
- Observe the outputs [on/off].
- Can be used as a building block for an embedding ML model.

But just how fast is it?



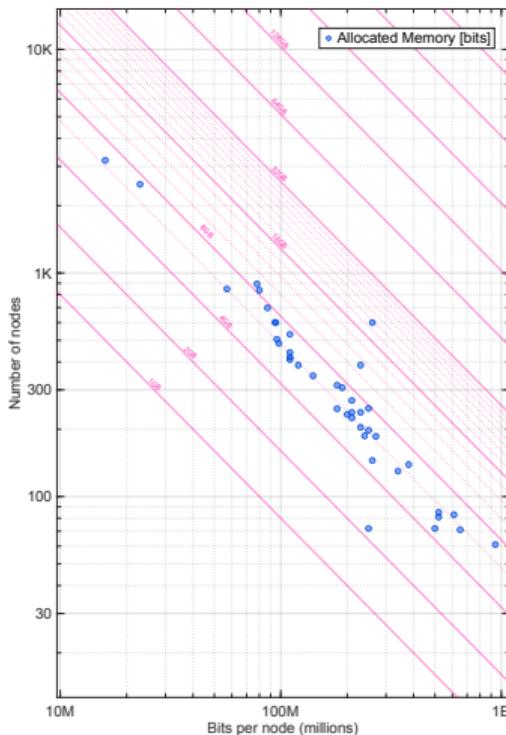


Eval Query on the Compiled Knowledge Graph

Eval Query Performance on GPUs:

- Latency: 200-300 ms per graph pass.
- Throughput: VRAM bound (see plot).
- Benchmarked on Nvidia GTX 1660 [6GB].
- Graph sizes: from ≈ 50 to ≈ 2000 nodes.
- Evals: from 16M to 1B per node per pass.

Q: Are these enough samples to estimate the Expectation Query?





Estimator for the Expected Value (i.e., Probability)

- A Boolean function $F(\mathbf{x})$ can be viewed as an indicator function: $F(\mathbf{x}) \in \{0, 1\}$.
- The event $\{F(\mathbf{X}) = 1\}$ has probability $\mathbb{E}[F(\mathbf{X})]$.
- **Monte Carlo estimator:**

$$\hat{P}_N = \frac{1}{N} \sum_{i=1}^N F(\mathbf{x}^{(i)}),$$

where each $\mathbf{x}^{(i)}$ is a random draw from the input distribution.

- By the Law of Large Numbers,

$$\lim_{N \rightarrow \infty} \hat{P}_N = \Pr[F(\mathbf{X}) = 1], \quad \text{almost surely.}$$

- Error decreases at rate $\mathcal{O}(1/\sqrt{N})$, analyzed via the Central Limit Theorem.



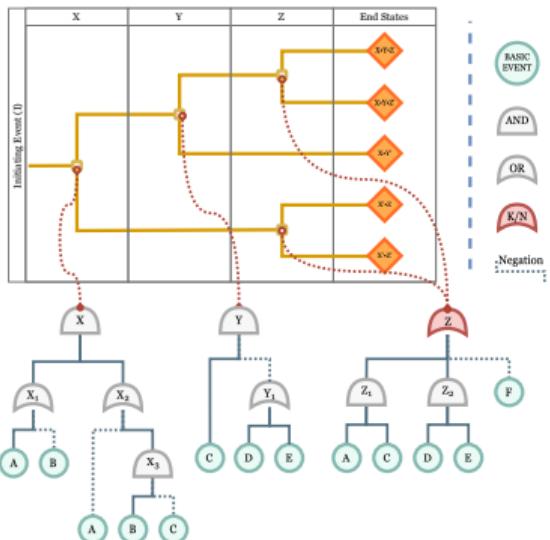
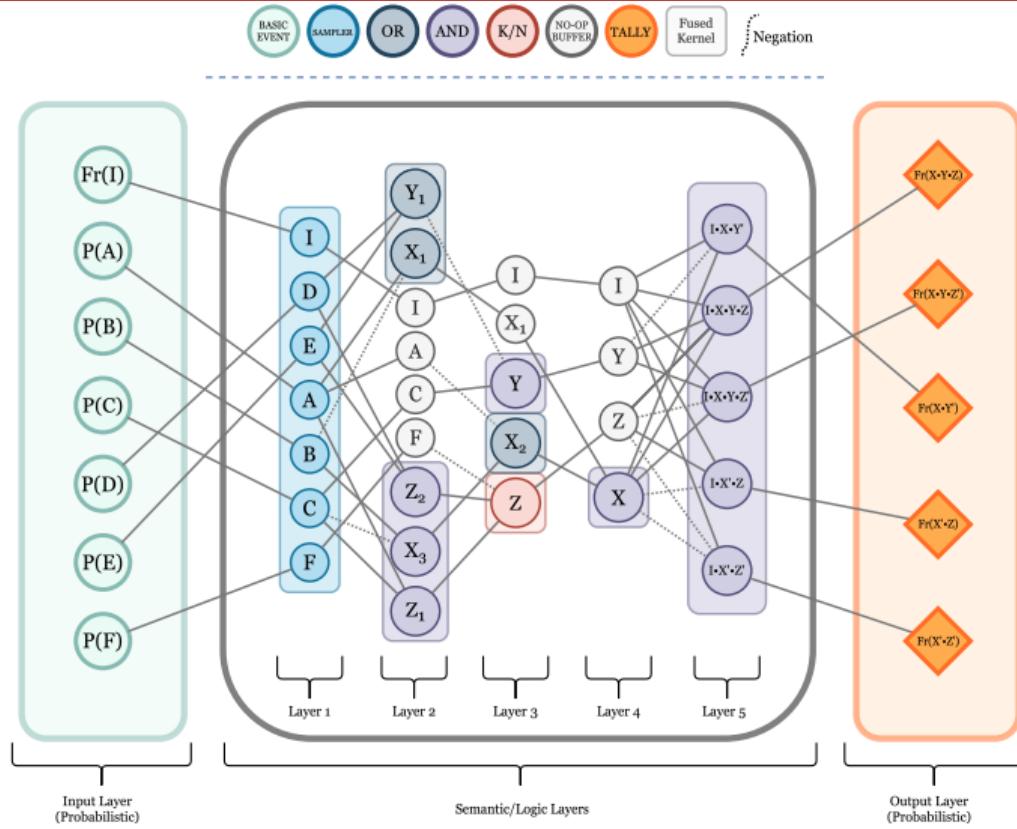
Monte Carlo Sampling

- Rather than summing or bounding all combinations of failures, *simulate* random draws of \mathbf{X} .
- Each Monte Carlo iteration:
 - 1 Sample $x_1, x_2, \dots, x_n \stackrel{\text{i.i.d.}}{\sim} \prod p(x_i)$.
 - 2 Evaluate the Boolean function $F(\mathbf{x})$ (cost is just logical gate evaluation).
 - 3 Collect whether $F(\mathbf{x}) = 1$ (failure) or 0 (success).
- Repeating for many samples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)}\}$ yields a *sample average* estimate of the probability.
- Benefits:
 - Bypasses explicit inclusion-exclusion expansions.
 - Straightforward to parallelize (evaluate each draw in separate threads or blocks).



└ PRA Models as Probabilistic Circuits

[Back to Working Example: One Initiating Event, Three Fault Trees, Six Basic Events, Five End States](#)





Preliminary Case Study



Overview: Aralia Dataset

- **Dataset Composition:** The Aralia collection consists of 43 trees.
- **Diverse Problem Sizes:** Small trees (e.g. 25–32 basic events) through mid-sized models with over 1,500 BEs.
- **Wide Probability Range:** Top-event probabilities spanning from rare events near 10^{-13} to fairly likely failures with probability above 0.7.
- **Model Variability:** Some trees are primarily AND/OR, others incorporate more advanced gates (K/N, XOR, NOT), providing thorough coverage of typical (and atypical) fault tree logic structures.



Benchmarking Setup: Hardware and Environment

■ Target Hardware:

- GPU: NVIDIA® GeForce GTX 1660 SUPER (6 GB GDDR6, 1,408 CUDA cores).
- CPU: Intel® Core™ i7-10700 (2.90 GHz, turbo-boost, hyperthreading).

■ Software Stack:

- SYCL-based (AdaptiveCpp/HipSYCL), with LLVM-IR JIT for kernel compilation.
- Compiler optimization at -O3 for efficient code generation.
- Repeated runs (5+) to mitigate transient variations.

■ Measured Time:

Includes entire wall-clock duration, from host-device transfers and JIT compilation to final result collection.



Monte Carlo Execution and Implementation

■ **Objective:** Compute TOP event probabilities for all 43 trees.

■ **Sampling Strategy:**

- Single pass per fault tree, generating as many samples as fit in 6 GB GPU memory.
- 128-bit Philox4x32x10 pseudo-random number generator, parallel threads.

■ **Bit-Packing Optimization:**

- Each group of 64 Monte Carlo outcomes stored in a single 64-bit word.
- Enables vectorized instructions (e.g. `popcount`) and reduces memory I/O.

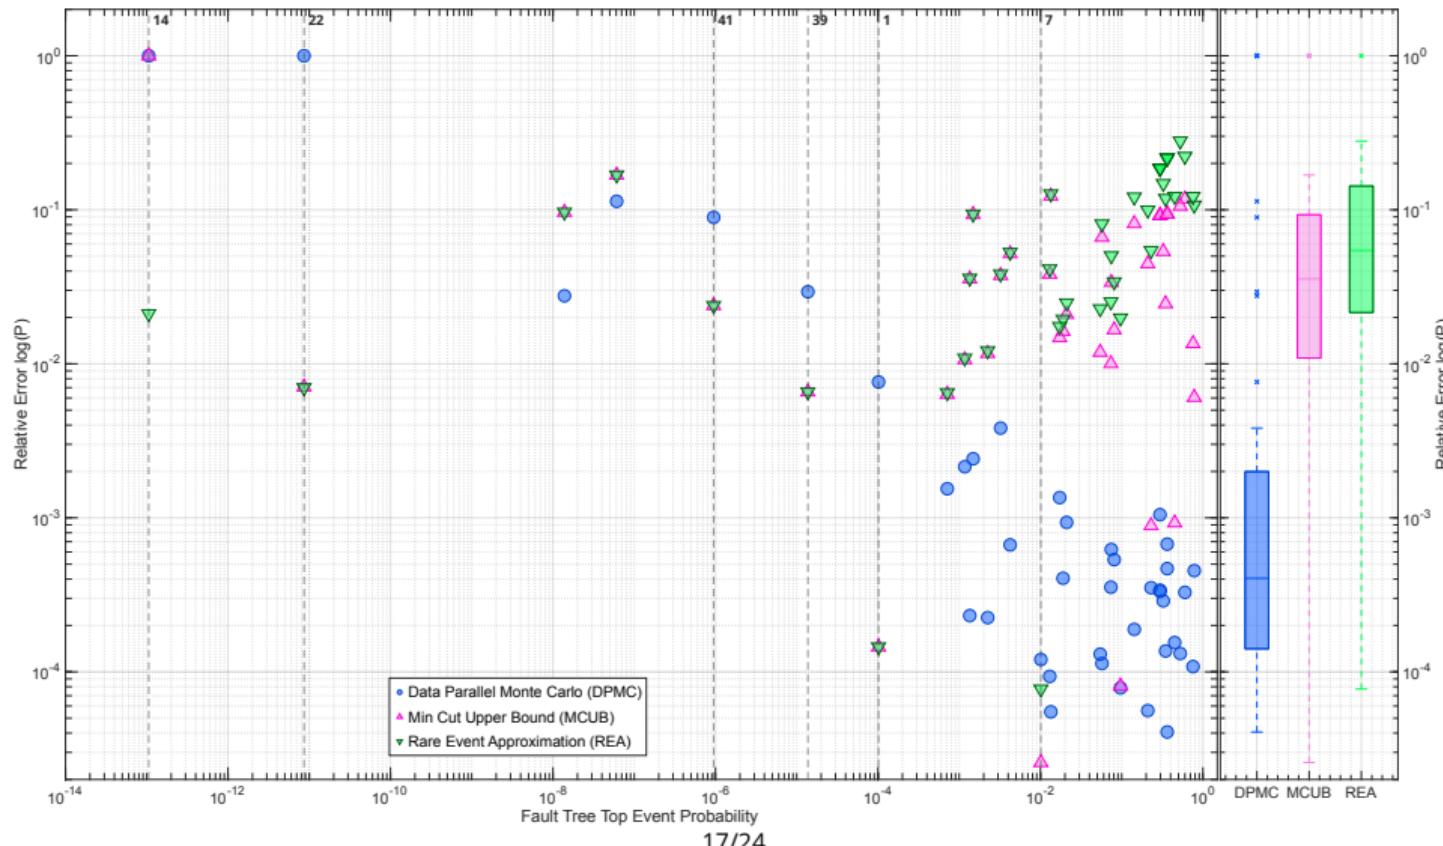
■ **Data Types:**

- Tallies in 64-bit integers.
- Probability accumulations in double precision (64-bit float).



└ Preliminary Case Study: Aralia Dataset

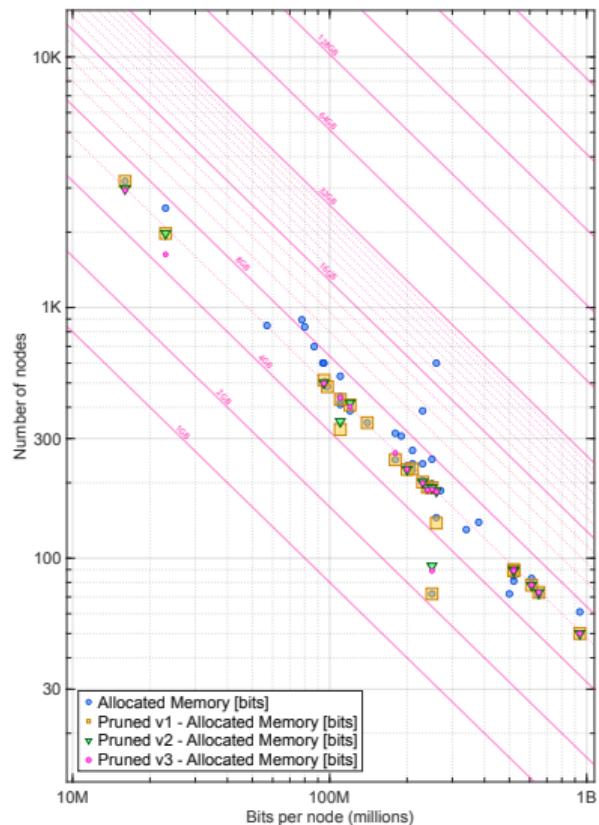
└ Accuracy Benchmark: Relative error (Log-probability), Data-Parallel Monte Carlo vs Min-Cut Upper Bound and Rare-Event Approximation



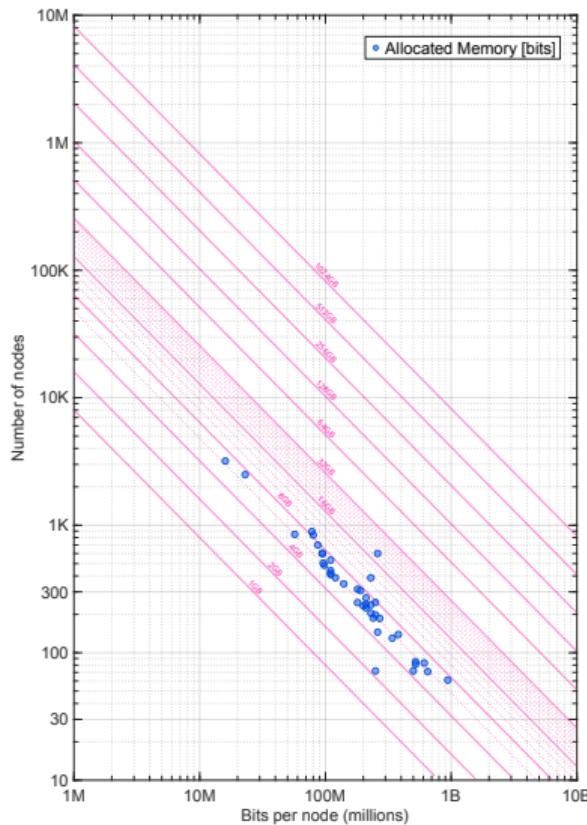


└ Preliminary Case Study: Aralia Dataset

└ Performance Benchmark (Memory Consumption): Sampled Bits Per Event Per Iteration



18/24





Limitations:

- Brute-force/naive Monte Carlo struggles when sampling rare-events.
 - Implement importance sampling: WIP.
- Brute-force/naive Monte Carlo is a poor strategy for sampling correlated events.

Next Steps:

- Benchmark on larger (G-PWR, G-MHTGR) models.

Future Work:

- Embedding models from Knowledge Graph, for semantic representation.
- Gradient computation on Knowledge Graph.
- Minimal cut set computation from Knowledge Graph.



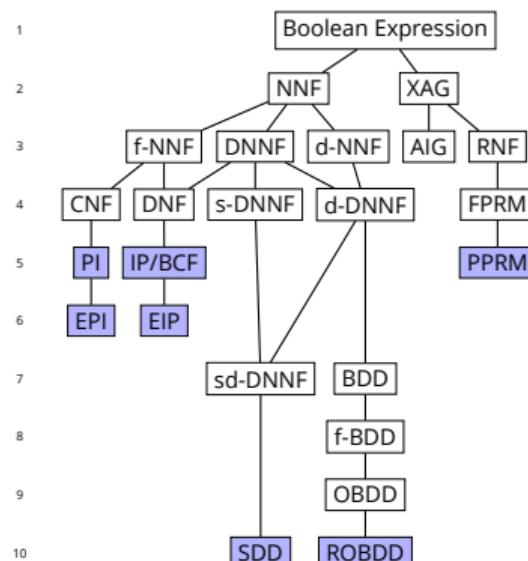
The End



└ Outlook

└ Knowledge Compilation

Hierarchy of compiled target languages. Blue nodes represent canonical forms.



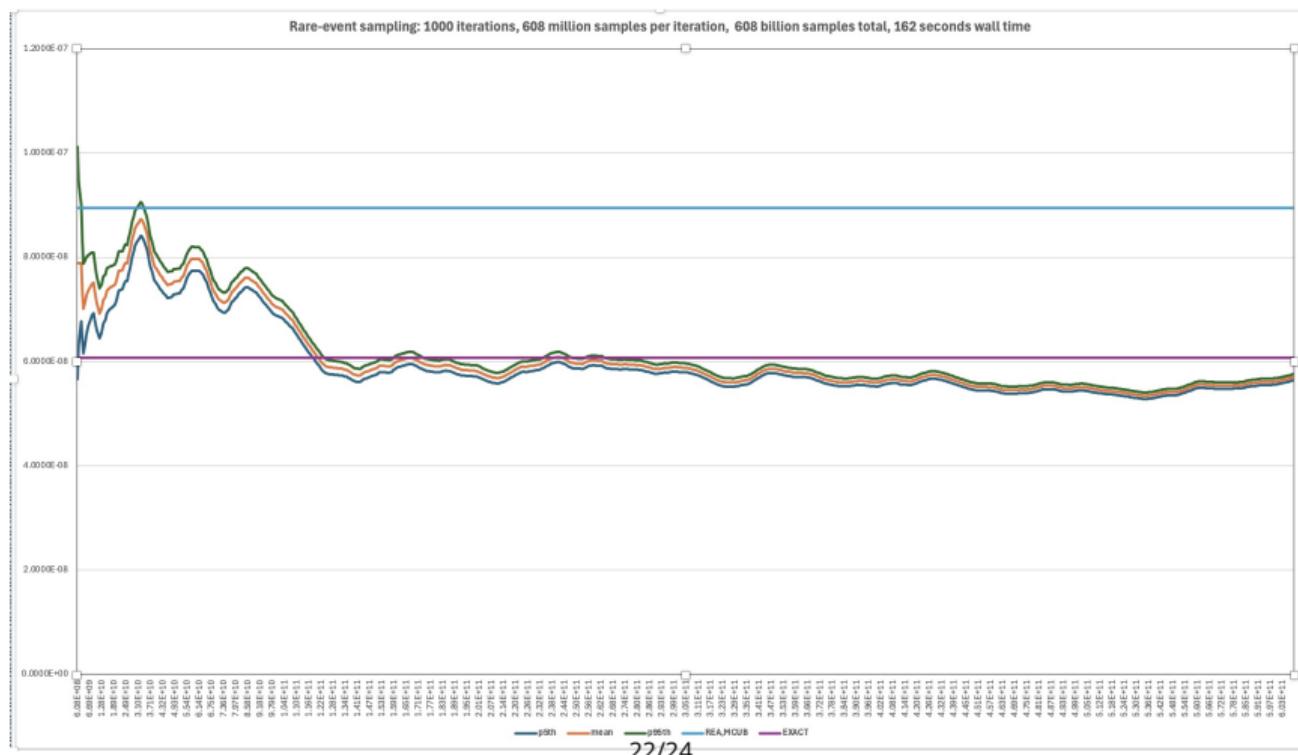
Acronym	Full form
NNF	Negation Normal Form
XAG	XOR-And-Inverter Graph
AIG	And-Inverter Graph
ANF/RNF	Algebraic/Ring Normal Form
f-NNF	Flat Negation Normal Form
DNNF	Decomposable Negation Normal Form
d-NNF	Deterministic Negation Normal Form
FPRM	Fixed Polarity Reed-Muller
CNF	Conjunctive Normal Form
DNF	Disjunctive Normal Form
s-DNNF	Smooth/Structured Decomposable Negation Normal Form
d-DNNF	Deterministic Decomposable Negation Normal Form
sd-DNNF	Smooth/Structured Deterministic Decomposable Negation Normal Form
PPRM	Positive Polarity Reed-Muller
PI	Prime Implicate
IP	Prime Implicant
BCF	Blake Canonical Form
EPI	Essential Prime Implicate
EIP	Essential Prime Implicant
BDD	Binary Decision Diagram
f-BDD	Free/Read-Once Binary Decision Diagram
OBDD	Ordered Binary Decision Diagram
SDD	Sentential Decision Diagram
RoBDD	Reduced Ordered Binary Decision Diagram



└ Preliminary Case Study: Aralia Dataset

└ Convergence Trends

Convergence over 1000 iterations, Aralia das9204





└ Preliminary Case Study: Aralia Dataset

└ Input Data

#	Fault Tree	Basic Events	Logic Gates					Minimal Cut Sets	Top Event Probability
			Total	AND	VOT	XOR	NOT		
1	baobab1	61	84	16	9	-	-	46,188	1.01708E-04
2	baobab2	32	40	5	6	-	-	4,805	7.13018E-04
3	baobab3	80	107	46	-	-	-	24,386	2.24117E-03
4	cea9601	186	201	69	8	-	30	130,281,976	1.48409E-03
5	chinese	25	36	13	-	-	-	392	1.17058E-03
6	das9201	122	82	19	-	-	-	14,217	1.34237E-02
7	das9202	49	36	10	-	-	-	27,778	1.01154E-02
8	das9203	51	30	1	-	-	-	16,200	1.34880E-03
9	das9204	53	30	12	-	-	-	16,704	6.07651E-08
10	das9205	51	20	2	-	-	-	17,280	1.38408E-08
11	das9206	121	112	21	-	-	-	19,518	2.29687E-01
12	das9207	276	324	59	-	-	-	25,988	3.46696E-01
13	das9208	103	145	33	-	-	-	8,060	1.30179E-02
14	das9209	109	73	18	-	-	-	8.20E+10	1.05800E-13
15	das9601	122	288	60	36	12	14	4,259	4.23440E-03
16	das9701	267	2,226	1,739	-	-	992	26,299,506	7.44694E-02
17	edf9201	183	132	12	-	-	-	579,720	3.24591E-01
18	edf9202	458	435	45	-	-	-	130,112	7.81302E-01
19	edf9203	362	475	117	-	-	-	20,807,446	5.99589E-01
20	edf9204	323	375	106	-	-	-	32,580,630	5.25374E-01
21	edf9205	165	142	30	-	-	-	21,308	2.09351E-01
22	edf9206	240	362	126	-	-	-	385,825,320	8.61500E-12



└ Preliminary Case Study: Aralia Dataset

└ Input Data

23	edfpa14b	311	290	70	-	-	-	105,955,422	2.95620E-01
24	edfpa14o	311	173	42	-	-	-	105,927,244	2.97057E-01
25	edfpa14p	124	101	42	-	-	-	415,500	8.07059E-02
26	edfpa14q	311	194	55	-	-	-	105,950,670	2.95905E-01
27	edfpa14r	106	132	55	-	-	-	380,412	2.09977E-02
28	edfpa15b	283	249	61	-	-	-	2,910,473	3.62737E-01
29	edfpa15o	283	138	33	-	-	-	2,906,753	3.62956E-01
30	edfpa15p	276	324	33	-	-	-	27,870	7.36302E-02
31	edfpa15q	283	158	45	-	-	-	2,910,473	3.62737E-01
32	edfpa15r	88	110	45	-	-	-	26,549	1.89750E-02
33	elf9601	145	242	97	-	-	-	151,348	9.66291E-02
34	ftr10	175	94	26	-	-	-	305	4.48677E-01
35	isp9601	143	104	25	1	-	-	276,785	5.71245E-02
36	isp9602	116	122	26	-	-	-	5,197,647	1.72447E-02
37	isp9603	91	95	37	-	-	-	3,434	3.23326E-03
38	isp9604	215	132	38	-	-	-	746,574	1.42751E-01
39	isp9605	32	40	8	6	-	-	5,630	1.37171E-05
40	isp9606	89	41	14	-	-	-	1,776	5.43174E-02
41	isp9607	74	65	23	-	-	-	150,436	9.49510E-07
42	jbd9601	533	315	71	-	-	-	150,436	7.55091E-01
43	nus9601	1,567	1,622	392	47	-	-	unknown	